


AOS Project

黄宏毅

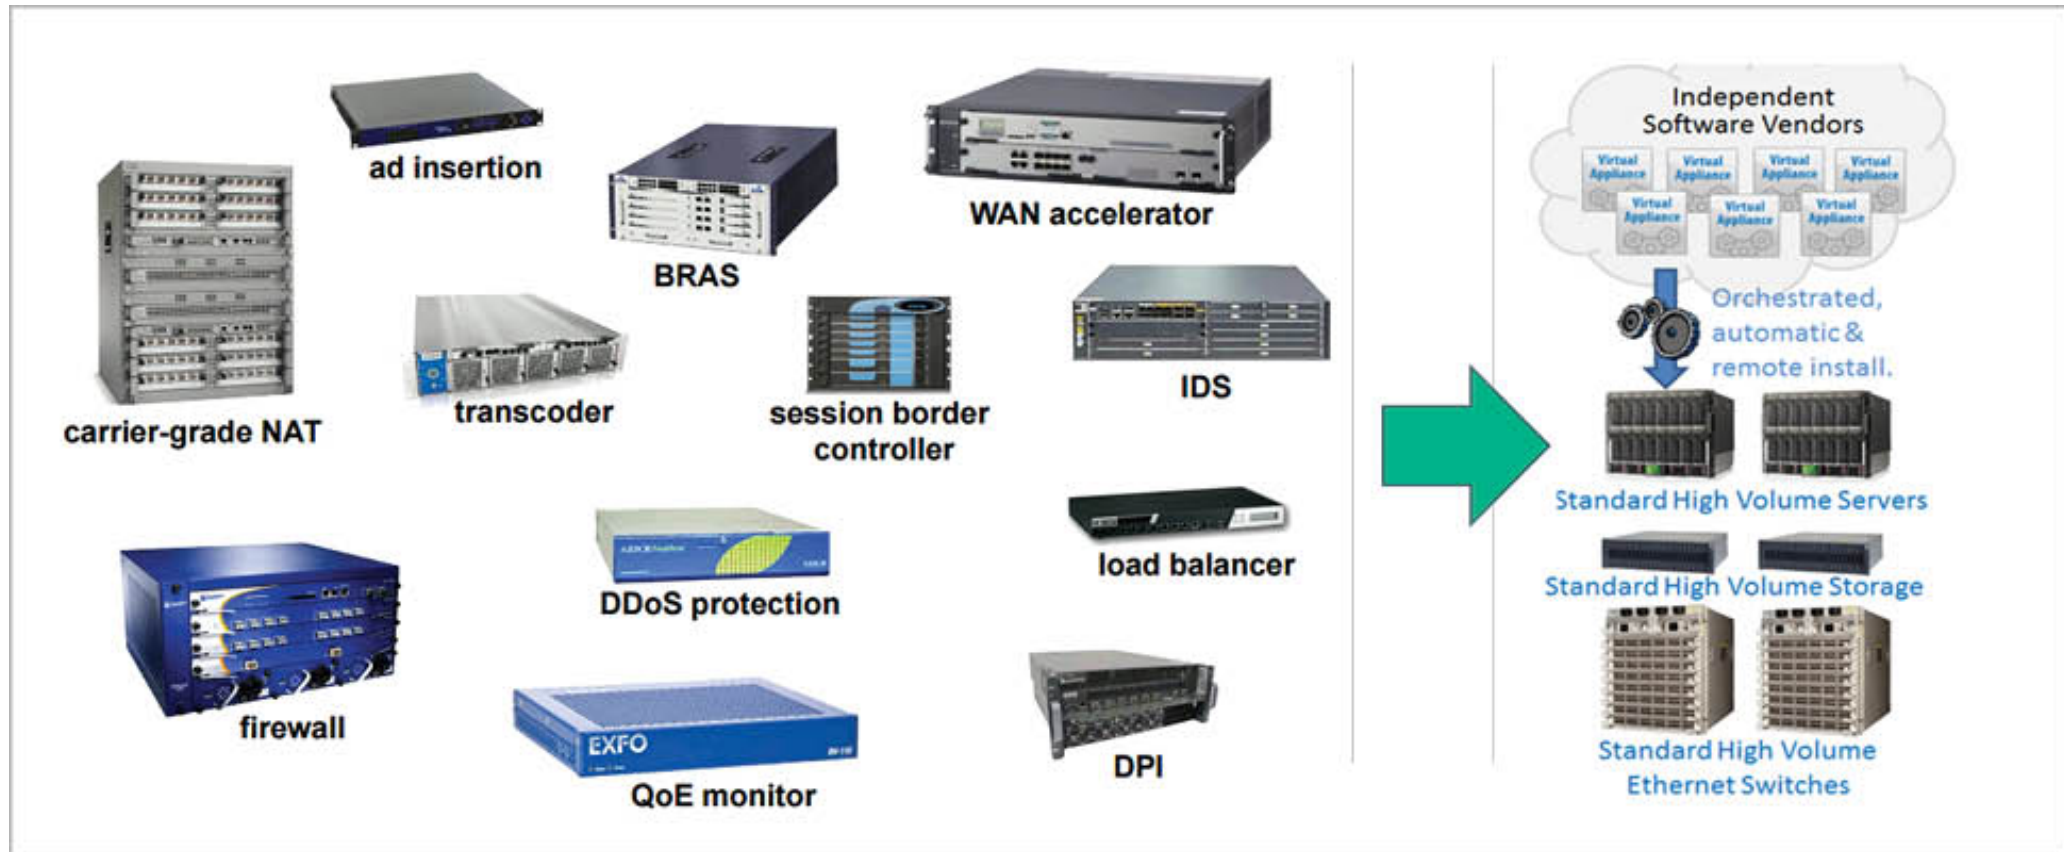
hhy17@mails.tsinghua.edu.cn

https://github.com/fjzzhongyi/aos_hw

SDN

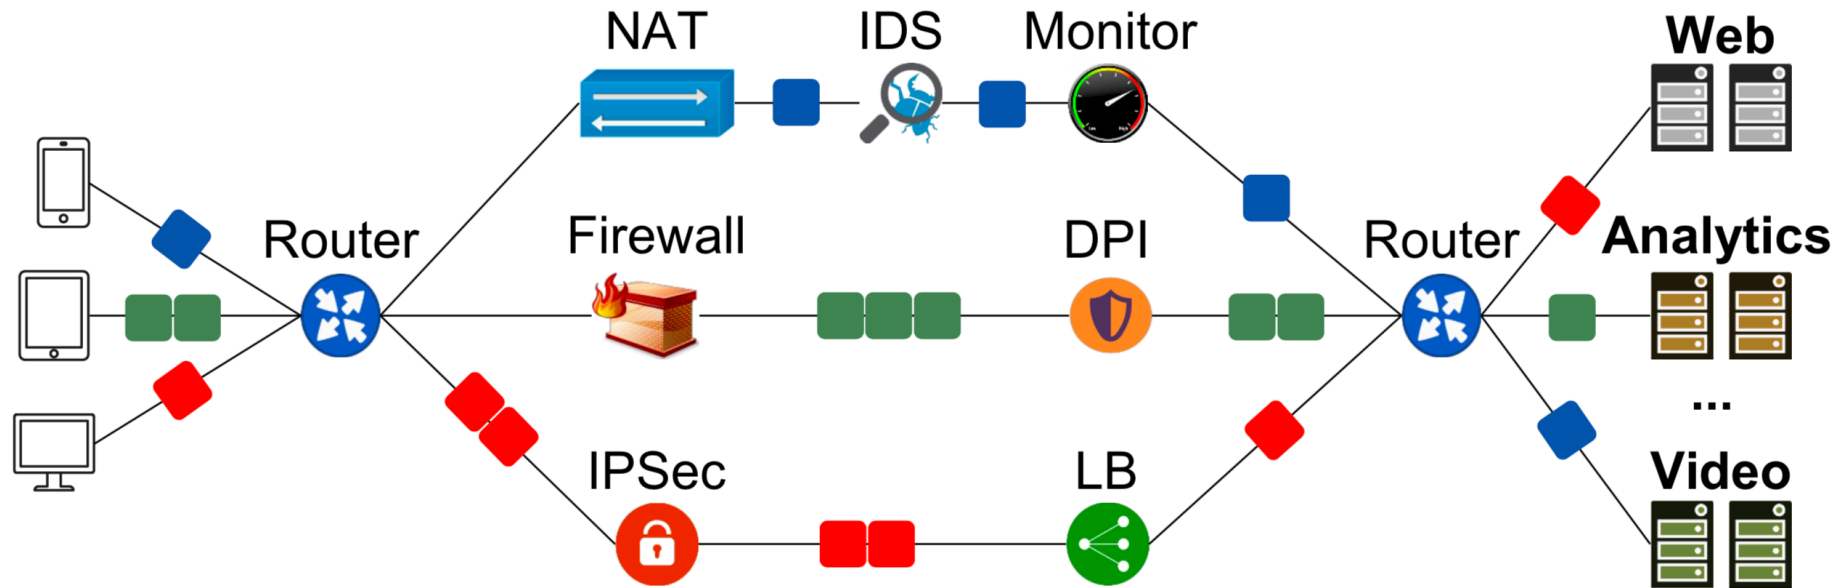


NFV

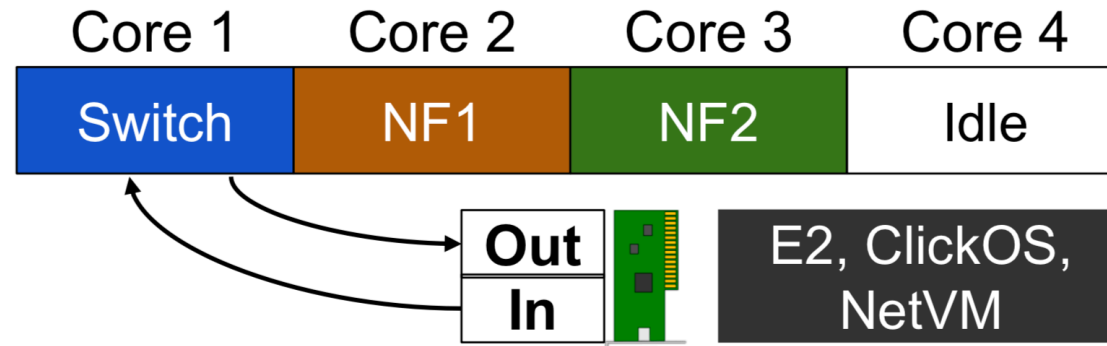


Metron: NFV Service Chains at the True Speed of the Underlying Hardware


- NF Chain



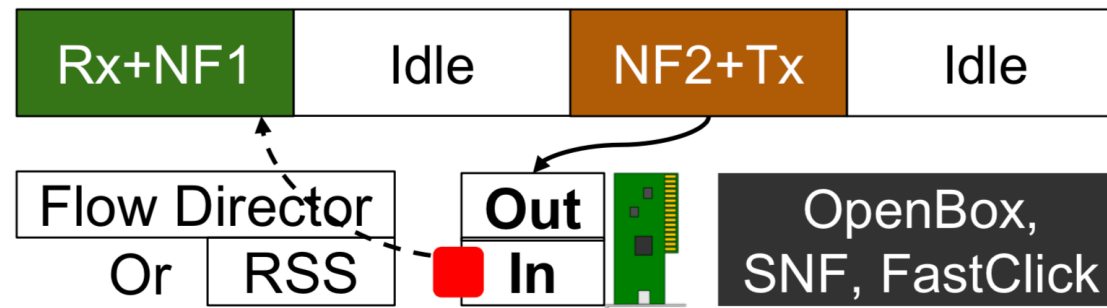
Dedicated switch



Pipeline

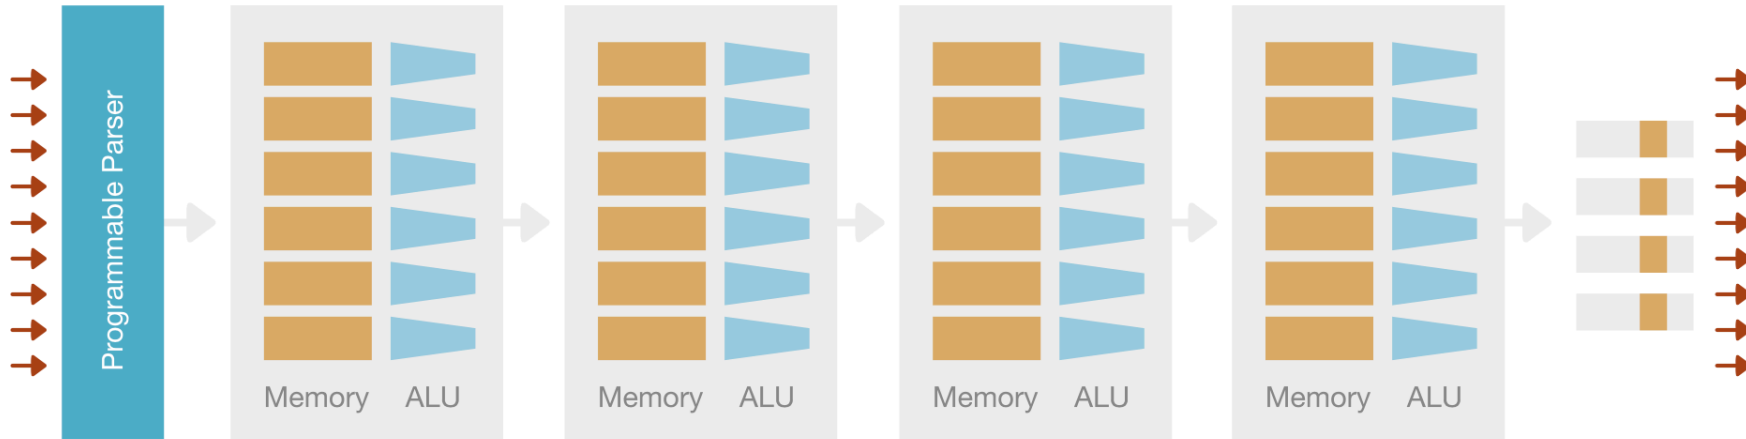


RSS/FD



Programmable Switch

match-action

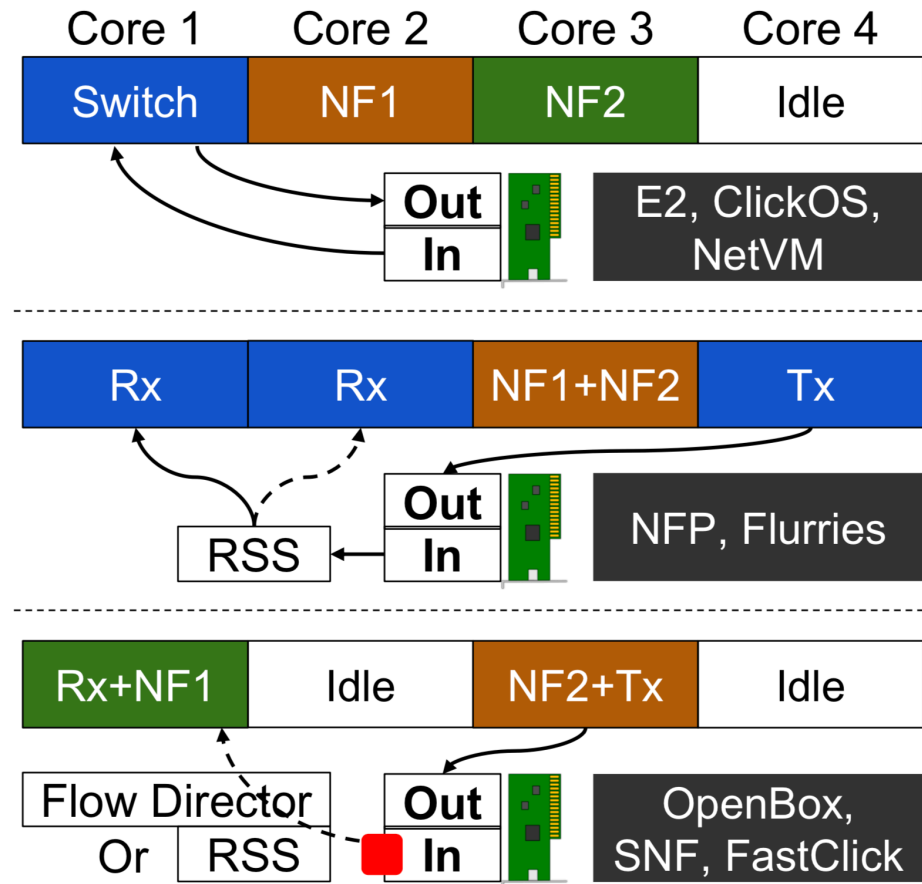


- Stateless read/write operation
- Tagging
- SDN → Dynamic tagging




Comments

- Globally shared state
 - Cross-core communication
 - Lock
- Scaling/Load Balancing
 - State migration




SafeBricks: Shielding Network Functions in the Cloud

- NetBricks
 - Rust
 - “Shared memory”
- Intel SGX
 - Enclave
 - IO optimization
- Compiler



ResQ: Enabling SLOs in Network Function Virtualization

- SLO: service level object
 - Predictable performance & latency
- Last-Level Cache(LLC)
 - Intel Cache Allocation Technology
 - Contention
- Solution: control DMA buffer



mTCP: A Highly Scalable User-level TCP Stack for Multicore Systems

- Linux TCP stack
 - 缺乏连接局部化
 - 共享的文件描述符空间
 - per-packet处理
 - 系统调用开销

mTCP

- User-level
- RSS
- Local socket
- DPDK
- Batch processing
- Lock-free queue (single consumer, single producer)

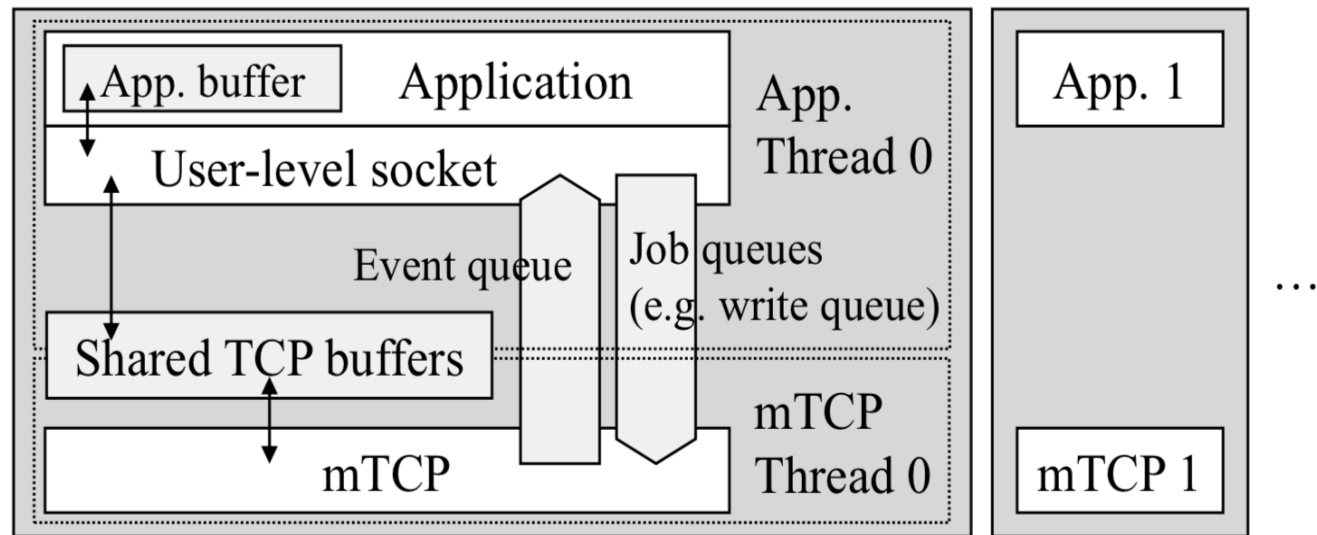


Figure 4: Thread model of mTCP.

Limitations:

1. 内核实现的协议/模块
2. 使用buffer会使copy开销大

FastSocket

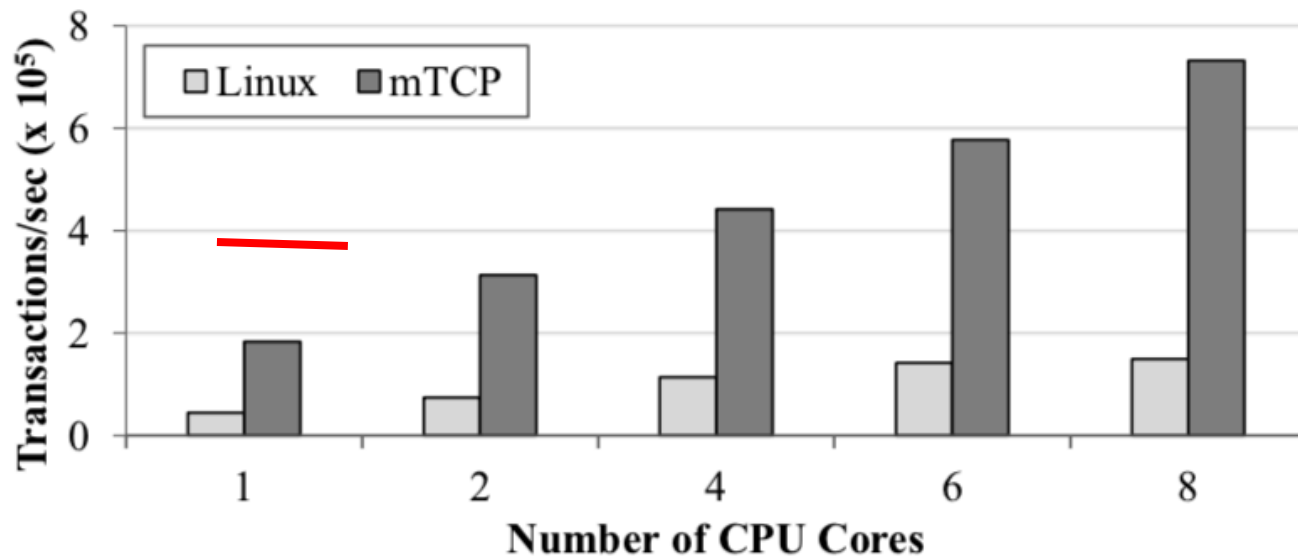
- BSD socket
 - Local listen table
 - Local established table
 - VFS
- RFD+hash
 - 流量不均衡

mTCP实验

- PSIO
 - Linux kernel 2.6.*
- DPDK
 - mTCP-dedicated DPDK
 - NIC support
- netmap

mTCP实验

- Config
 - VMware
 - e1000 virtual NIC
 - Ubuntu 14/16
 - LAN
- Setup
 - Hugepage
 - igb module
 - Configure NIC
 - no ping-pong
 - wget 64B



$t=0$

Figure 10: Performance of ab as a function of the number of cores. The file size is 64B and 8K concurrent connections are used.

HyperKernel实验

- Ubuntu16.04, 18.04
 - make hv6-verify failed!
- Docker
 - spec.py → sys_close

```
# decrement file refcnt  
new.files[fn].refcnt[(pid, fd)]
```

```
=====
FAIL: test_sys_close (__main__.HV6)
-----
Traceback (most recent call last):
  File "o.x86_64/hv6/hv6py/kernel/spec/main.py", line 410, in <lambda>
    self._syscall_generic(syscall)
  File "o.x86_64/hv6/hv6py/kernel/spec/main.py", line 545, in _syscall_generic
    cond, newstate = getattr(spec, name)(self.state, *args)
  File "/home/hv6/o.x86_64/hv6/hv6py/kernel/spec/spec/specs.py", line 794, in sys_close
    new.files[fn].refcnt[(pid, fd)] -= 0
  File "/home/hv6/o.x86_64/hv6/hv6py/base.py", line 60, in __sub__
    return self._imap.__isub__(self._instance, value)
  File "/home/hv6/o.x86_64/hv6/hv6py/base.py", line 115, in __isub__
    return self.isub(instance._fields[self._name], instance._path, value)
  File "/home/hv6/o.x86_64/hv6/hv6py/base.py", line 627, in isub
    assert value == 1
AssertionError

-----
Ran 1 test in 8.456s

FAILED (failures=1)
hv6/Makefrag:169: recipe for target 'hv6-verify' failed
make: *** [hv6-verify] Error 1
```